



NICIS

NATIONAL INTEGRATED
CYBERINFRASTRUCTURE SYSTEM

DIRISA

Introduction to Data Visualization

Minenhle Mabuyakhulu

Rosemary Erawemen

24 June 2025

AN INITIATIVE OF:



science, technology
& innovation

Department:
Science, Technology and Innovation
REPUBLIC OF SOUTH AFRICA



CSIR
Touching lives through innovation

80th
anniversary



Introduction

- Visualization can be used to:
- Better scrutinize a dataset (e.g., identifying outliers);
- Obtain a better understanding of data (breakdown);
- Find relationships and patterns in the dataset (pattern recognition)
- Sharing results and concepts with others (communication).

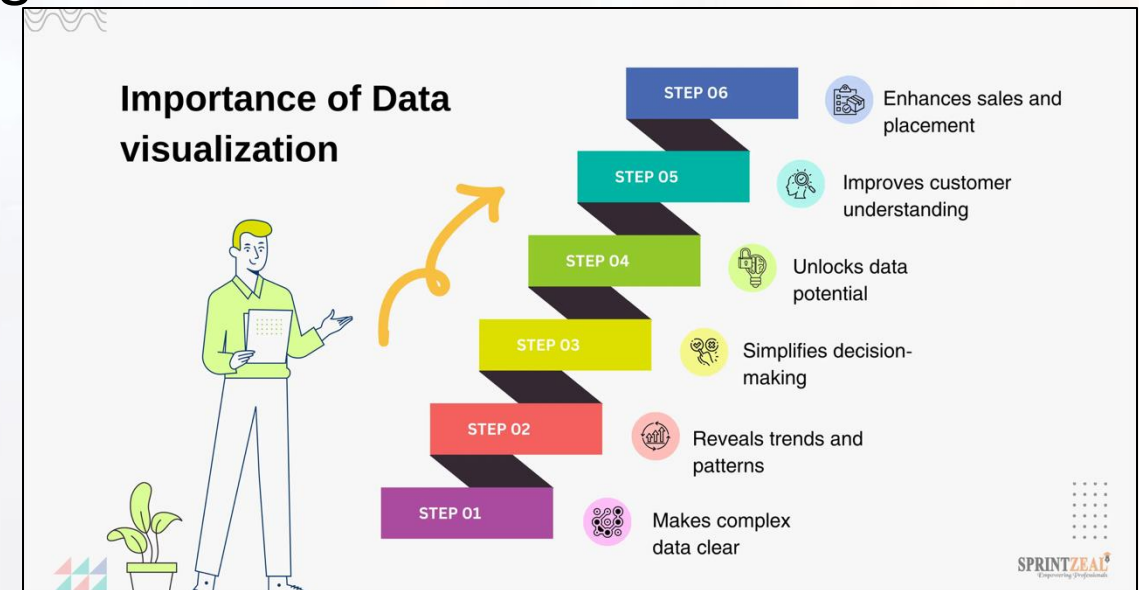


What is Data Visualization?

- It's a way to **turn numbers and data into pictures.**
- These pictures can be **charts, graphs, or maps.**
- It helps us **see and understand data more easily.**
- Makes complex data **look simple and clear.**
- Helps us to **find patterns, compare things, and tell stories with data.**

Goals of Data Visualization

- To make data easy to understand
- To find patterns or trends in the data
- To spot mistakes or problems quickly
- To share results clearly with others
- To help in decision-making



Source: [1]



Types of Visualization

- In python ***matplotlib*** or ***seaborn*** can be used.
- Some examples of different visualization methods are:
 - **Line Graph**
 - **Scatterplot**
 - **Histogram**
 - **Heat Map**
 - **Bar And Stacked Graph**
 - **Pie Chart**
 - **Violin And Swarm Plot**

Common Graphs:

- **Line Graph** – shows trends over time (e.g., sales over months)
- **Scatter Plot** – shows relationships between two things (e.g., height vs. weight)
- **Histogram** – shows how often values appear (e.g., test scores)
- **Heatmap** – uses color to show patterns (e.g., correlation between variables)
- **Bar Chart / Stacked Bar** – compares categories (e.g., number of students per course)
- **Pie Chart** – shows parts of a whole (e.g., percentage of budget spent)
- **Violin Plot / Swarm Plot** – shows how data is spread out and grouped

The background features a dark blue field with vertical columns of binary code (0s and 1s) in a light blue, monospace font. Interspersed among the code are numerous out-of-focus circular bokeh lights in shades of blue and orange, creating a digital, high-tech atmosphere.

01

Understanding Data before Visualization



Purpose of Reading and Understanding the Dataset for Visualization

- Know what the data is about
 - Understand the topic and context.
- Understand each column or number
 - For example, "Sales," "Date," or "Product" — this helps choose the right chart.
- Check if the data fits visualization
 - Look at size, type (numerical or categorical), and if there are enough values.
- Identify the insight or message
 - Decide if you want to show trends, comparisons, distributions, or relationships.



Purpose of Reading and Understanding the Dataset for Visualization

- Make sure data is organized
 - It should be clear and ready for tools like Excel, Python, or Power BI.
- Spot basic issues
 - Watch out for missing values, outliers, or mislabeled columns to avoid confusing visuals.
- Select relevant data only
 - Focus on columns and values that support your story; exclude unnecessary data.



Dataset

- Small seaborn online datasets can be obtained to test some visualization techniques.

```
1 import seaborn as sns
2 sns.get_dataset_names()

['anagrams',
 'anscombe',
 'attention',
 'brain_networks',
 'car_crashes',
 'diamonds',
 'dots',
 'exercise',
 'flights',
 'fmri',
 'gammas',
 'geyser',
 'iris',
 'mpg',
 'penguins',
 'planets',
 'taxis',
 'tips',
 'titanic']
```

```
1 Data = sns.load_dataset("titanic")
2 display(Data)
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True
...
886	0	2	male	27.0	0	0	13.0000	S	Second	man	True	NaN	Southampton	no	True
887	1	1	female	19.0	0	0	30.0000	S	First	woman	False	B	Southampton	yes	True
888	0	3	female	NaN	1	2	23.4500	S	Third	woman	False	NaN	Southampton	no	False
889	1	1	male	26.0	0	0	30.0000	C	First	man	True	C	Cherbourg	yes	True
890	0	3	male	32.0	0	0	7.7500	Q	Third	man	True	NaN	Queenstown	no	True

The background of the slide features a dark blue field with vertical columns of binary code (0s and 1s) in a light blue, monospace font. Interspersed among the code are numerous out-of-focus circular bokeh lights in shades of blue and orange, creating a digital, high-tech atmosphere.

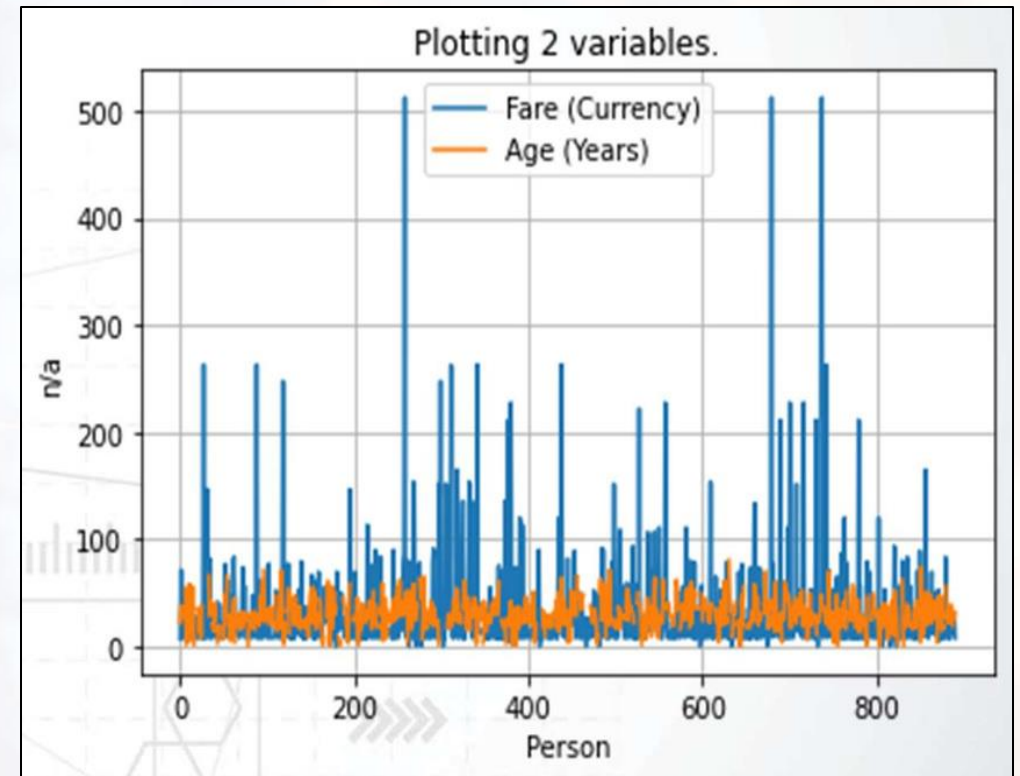
02

Common types of Graphs and When to Use Them

Line Graph

- Understanding of sequential such as time or the difference between similar variables.
 - It's used when you want to **see a trend or change**
 - Helps us see the **direction** and **speed** of change.

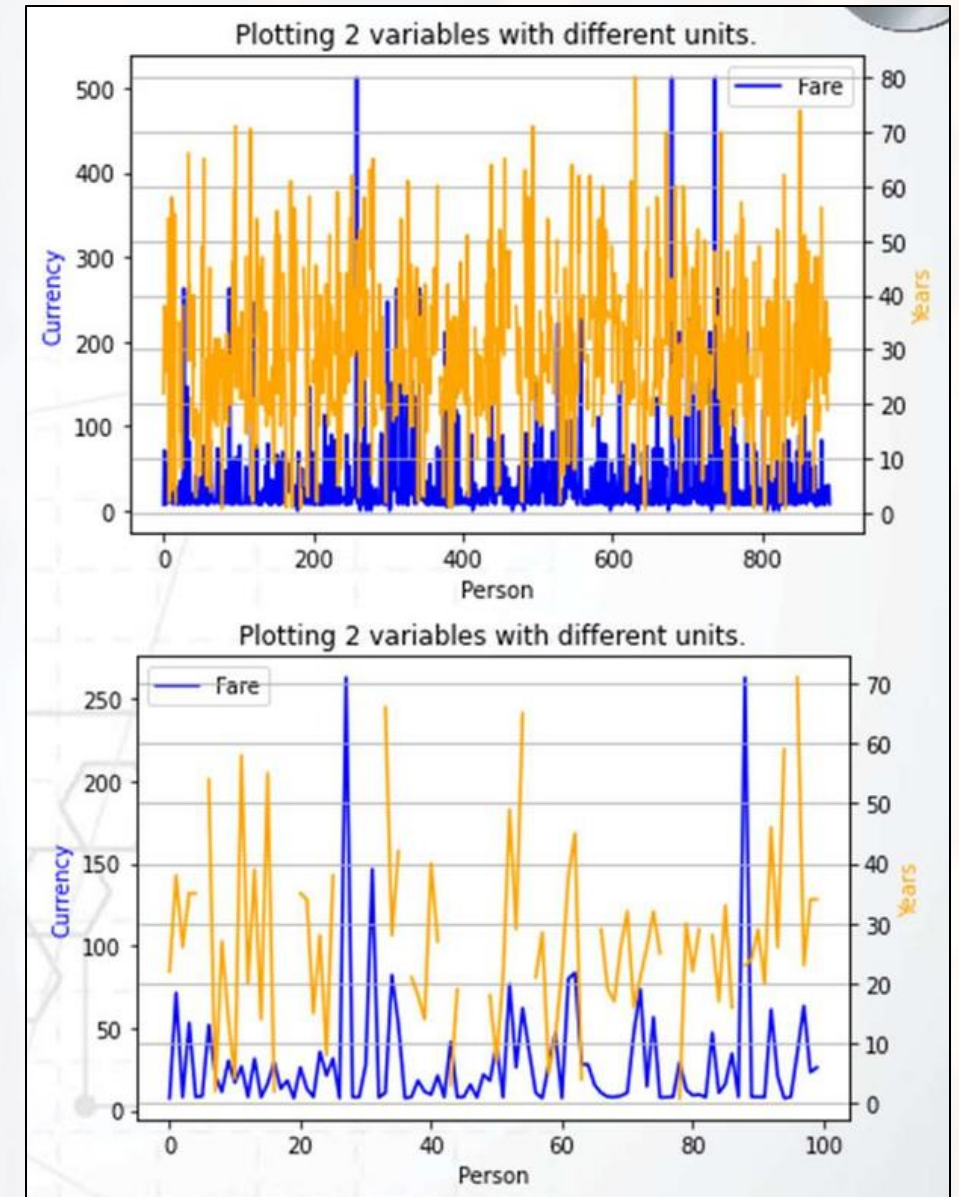
```
1 import matplotlib.pyplot as plt
2
3 plt.plot(Data['fare'],label='Fare (Currency)')
4 plt.plot(Data['age'],label='Age (Years)')
5 plt.grid()
6 plt.xlabel('Person')
7 plt.ylabel('n/a')
8 plt.title('Plotting 2 variables.')
9 plt.legend()
```



Line Graph

- In situations where you want to visually compare multiple variables with the same x values, but vastly different y ranges, a different y axis can be used for each variable.


```
1 fig, ax = plt.subplots()
2 ## FIRST PLOT
3 ax.plot(Data['fare'], color='blue', label = 'Fare (Currency)')
4 # set x-axis label
5 ax.set_xlabel('Person')
6 # set y-axis label
7 ax.set_ylabel('Currency', color='blue')
8 ## SECOND PLOT
9 ax2 = ax.twinx()
10 ax2.plot(Data['age'], color='orange', label = 'Age (Years)')
11 ax2.set_ylabel('Years', color='orange')
12 ax2.set_title('Plotting 2 variables with different units')
13 ax2.grid()
```



Scatterplot

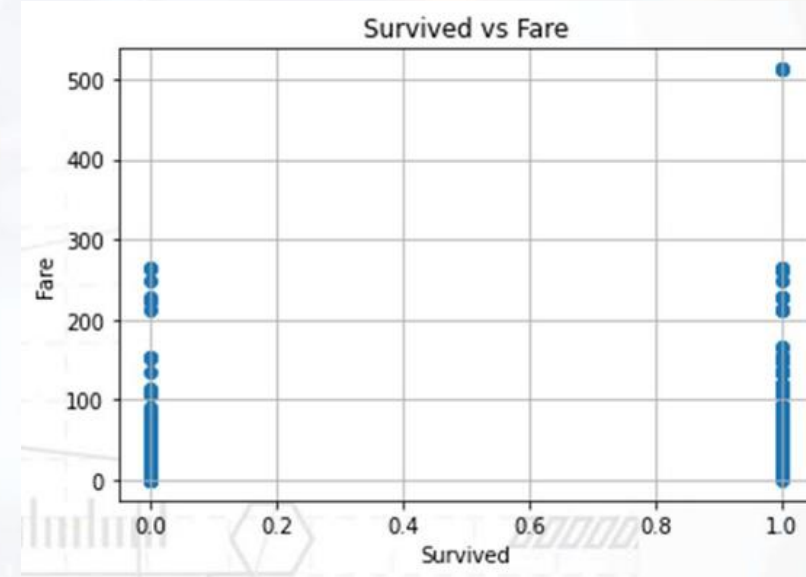

- A Scatterplot is useful to visualize the relationship between 2 variables and show how 2 variables might be related.
- Both datasets must have the same resolution so that a 1-on-1 comparison can be done.
- Scatter plots help us understand **correlation** — how closely two things move together

```
1 plt.scatter(Data['age'],Data['fare'])  
2 plt.xlabel('Age')  
3 plt.ylabel('Fare')  
4 plt.title('Age vs Fare')  
5 plt.grid()
```



Scatterplot

- Comparison: It is important to understand the meaning of the data plotted




- In this example both graphs show no clear relationship between the 2 variable.

Histogram


- A histogram displays a frequency distribution of a variable, and this can be used to determine if there are any 'intensity' ranges that appears more (or less) in the sample set.
- To show the **distribution** of data and also to group data into ranges and count how many values are in each.
- It helps us understand the **shape of the data** — like whether it's balanced, or has more values on the high or low side.

```
1 plt.hist(Data['age'], bins=30, rwidth=0.95) #,density=True
2 plt.ylabel('Nuber of samples within\nsample value range.')
3 plt.xlabel('Age')
4 plt.title('Age distribution')
5 #plt.grid()
```



Histogram

- Seaborn can also be used for plotting and is in many cases simpler to use than matplotlib.
- The number of bins (distribution resolution) can be set using 'bins' and the 'kde' (kernel density [probability] estimate) can be calculated and added to the graph using `kde = True`.




```
1 import seaborn as sns
2 sns.set_theme()
3 p = sns.displot(Data['age'], bins=30, kde=True)
4 p.set(xlabel='Age range', ylabel='Nuber of samples within value range.', title='Age dist.')
```



Histogram

- Seaborn can also plot dataframes directly (multiple variables) without any conversion needed.




```
1 sns.set_theme()
2 p = sns.displot(data = Data, x = 'age', hue = 'sex', bins = 30)
3 p.set(xlabel = 'Age Binned',
4       ylabel = 'Number of samples falling in age bin.',
5       title = 'Distribution')
```

Bar Graph

- We use bar graphs to **compare things side by side** — like how much different neighborhoods earn, or how many homes were sold in each city.
- Bar graphs make it easy to see **which group is the biggest or smallest**.
- You can also see patterns — like if one group is far ahead of the rest, or if values are close together. It helps make quick decisions or spot trends in simple data.

```
df['ocean_proximity'].value_counts().plot(kind='bar',  
                                           figsize=(8, 5), color='orange')  
plt.xlabel("Ocean Proximity")  
plt.ylabel("Count")  
plt.title("Number of Households Near Ocean")  
plt.show()
```



Bar Graphs - Stacked

- This can be used to display categorical data eg. When there are several variables with multiple attributes that the user wants to compare.

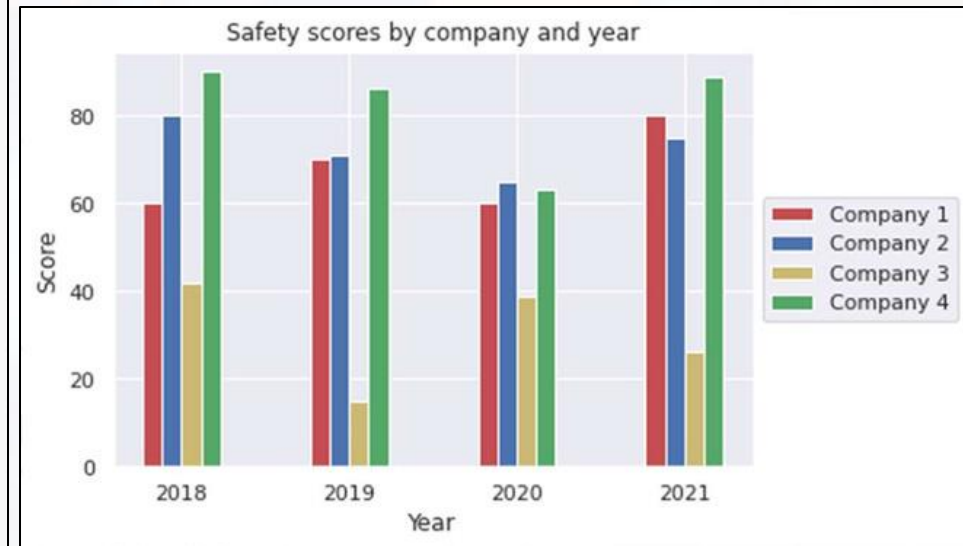


```
8 ind = np.array([0, 3, 6, 9]) # the x locations for the groups
9 width = 0.35 # the width of the bars
10
11 fig = plt.figure()
12 ax = fig.add_subplot(111)
13 p1 = ax.bar(ind, c1, width, color='r') #, yerr=menStd
14 p2 = ax.bar(ind+width, c2, width, color='b')
15 p3 = ax.bar(ind+width*2, c3, width, color='y')
16 p4 = ax.bar(ind+width*3, c4, width, color='g')
17
18 ax.set_xlabel('Year')
19 ax.set_ylabel('Score')
20 ax.set_title('Safety scores by company and year')
21 ax.set_xticks(ind+width*1.5)
22 ax.set_xticklabels(('2018', '2019', '2020', '2021'))
23
24 ax.legend(('Company 1', 'Company 2', 'Company 3', 'Company 4'),
25          loc='center left', bbox_to_anchor=(1, 0.5))
```


Bar Graphs – Side by Side

- There are different options for how to display a bar graph e.g., stacked, on the side axis, side-by-side, etc.

```
8 ind = np.array([0, 3, 6, 9]) # the x locations for the groups
9 width = 0.35 # the width of the bars
10
11 fig = plt.figure()
12 ax = fig.add_subplot(111)
13 p1 = ax.bar(ind, c1, width, color='r') #, yerr=menStd
14 p2 = ax.bar(ind+width, c2, width, color='b')
15 p3 = ax.bar(ind+width*2, c3, width, color='y')
16 p4 = ax.bar(ind+width*3, c4, width, color='g')
17
18 ax.set_xlabel('Year')
19 ax.set_ylabel('Score')
20 ax.set_title('Safety scores by company and year')
21 ax.set_xticks(ind+width*1.5)
22 ax.set_xticklabels(('2018', '2019', '2020', '2021'))
23
24 ax.legend(('Company 1', 'Company 2', 'Company 3', 'Company 4'),
25          loc='center left', bbox_to_anchor=(1, 0.5))
```



Bar Graph – Side by Side

- Sometimes (depending on the input data format) seaborn is not the easiest option even if, at first glance, it looks like it...
- Even if it looks simple, this took more time to figure out what the data format should be than plotting the previous example.



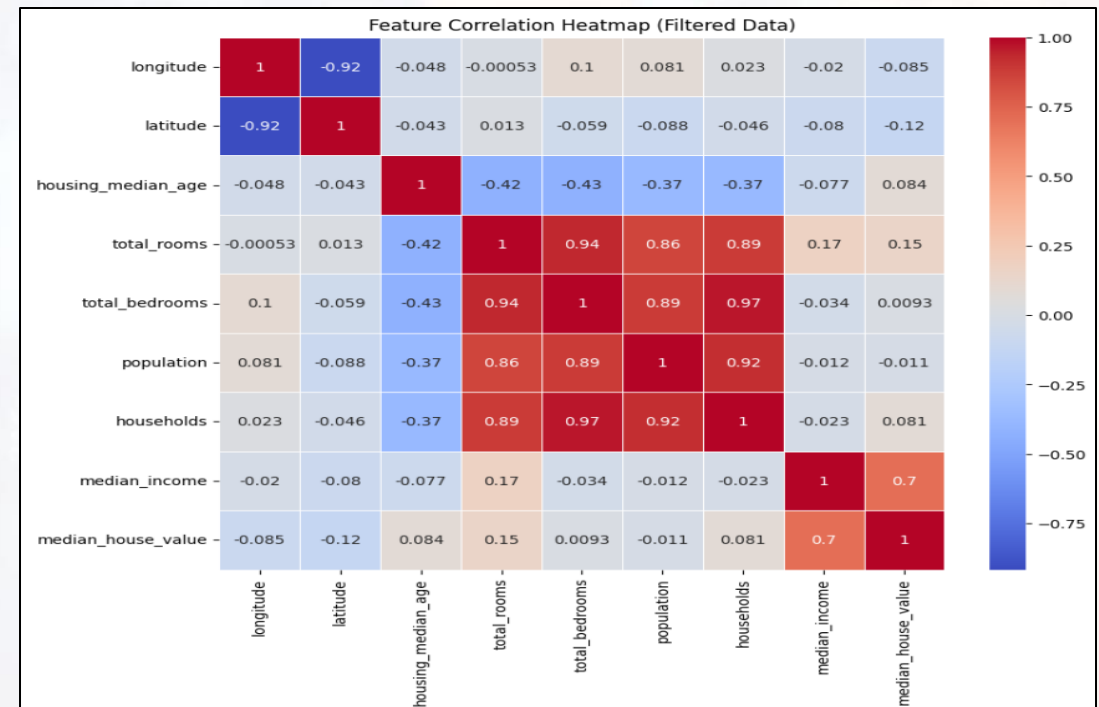


Heatmap

- A heatmap allows for data to be visualized in 2 dimensions with the magnitude indicated by the colour.
 - Darker or brighter colors mean higher or lower numbers.
- It helps show patterns in large amounts of data.

```
[31] # Select only numeric columns for correlation
numeric_df = sampled_df.select_dtypes(include=['float64', 'int64'])

plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
plt.title("Feature Correlation Heatmap (Filtered Data)")
plt.show()
```





Heatmap

- Another method to get the dataframe data in the needed format is to use 'groupby' and 'unstack' (count the number of instances).

	survived	pclass	sex	age	...	deck	embark_town	alive	alone
0	0	3	male	22.0	...	NaN	Southampton	no	False
1	1	1	female	38.0	...	C	Cherbourg	yes	False
2	1	3	female	26.0	...	NaN	Southampton	yes	True
3	1	1	female	35.0	...	C	Southampton	yes	False
4	0	3	male	35.0	...	NaN	Southampton	no	True
..
886	0	2	male	27.0	...	NaN	Southampton	no	True
887	1	1	female	19.0	...	B	Southampton	yes	True
888	0	3	female	NaN	...	NaN	Southampton	no	False
889	1	1	male	26.0	...	C	Cherbourg	yes	True
890	0	3	male	32.0	...	NaN	Queenstown	no	True

```
3 print(df)
4 df_m = data.groupby(["sex", "alive"]).size().unstack(level=0)
5 display(df_m)
```

sex	female	male
alive		
no	81	468
yes	233	109

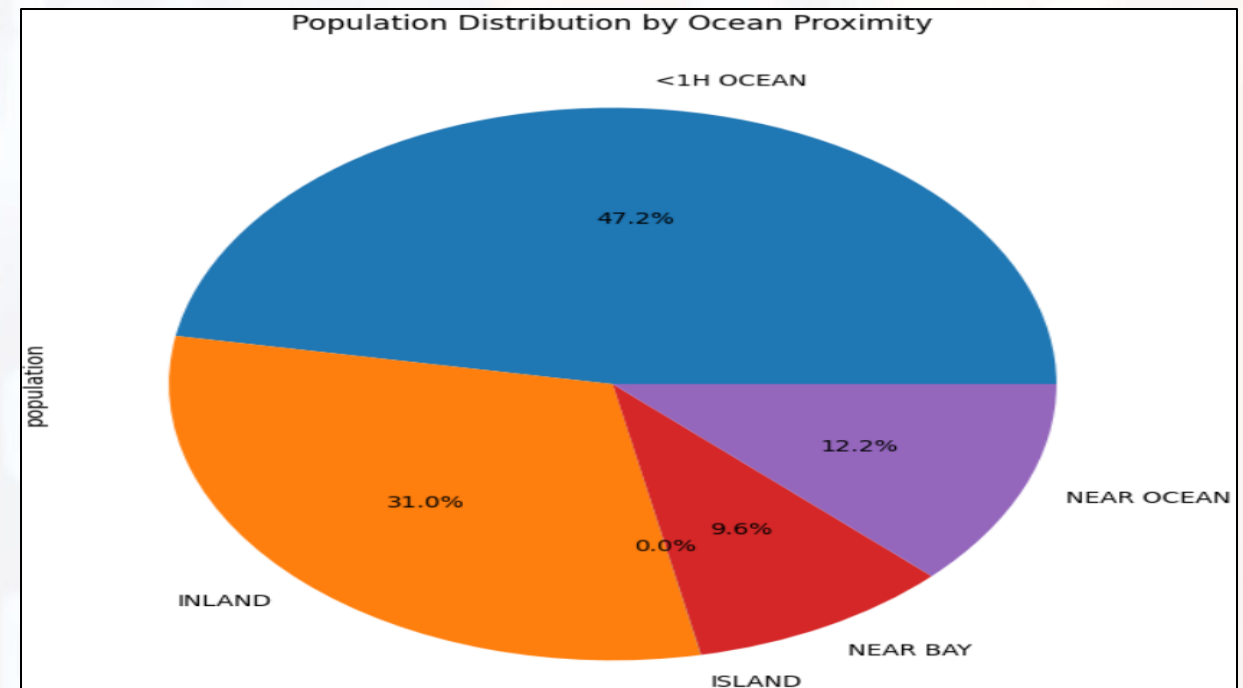
```
sns.heatmap(df_m, annot=True)
```



Pie Chart

- Used to show the share of the total for each variable or class.
- Each slice represents a percentage or part of a group.
- Helps see how different parts compare to each other.

```
df.groupby('ocean_proximity')['population'].sum().plot(kind='pie',  
                                                    autopct='%1.1f%%', figsize=(8, 8))  
plt.title("Population Distribution by Ocean Proximity")  
plt.show()
```







Pie Char

- Similar to heatmap, 'groupby' and 'unstack' can be used to count the number of instances.

```
6 plot = datapie.plot.pie(y='2018', figsize=(5, 5))  
7 plt.legend(bbox_to_anchor=(0.9,0.5), loc="center left")
```



```
1 datapie = datapie.transpose()  
2  
3 plot = datapie.plot.pie(y='company 1', figsize=(5, 5))  
4 plt.legend(bbox_to_anchor=(0.9,0.5), loc="center left")
```


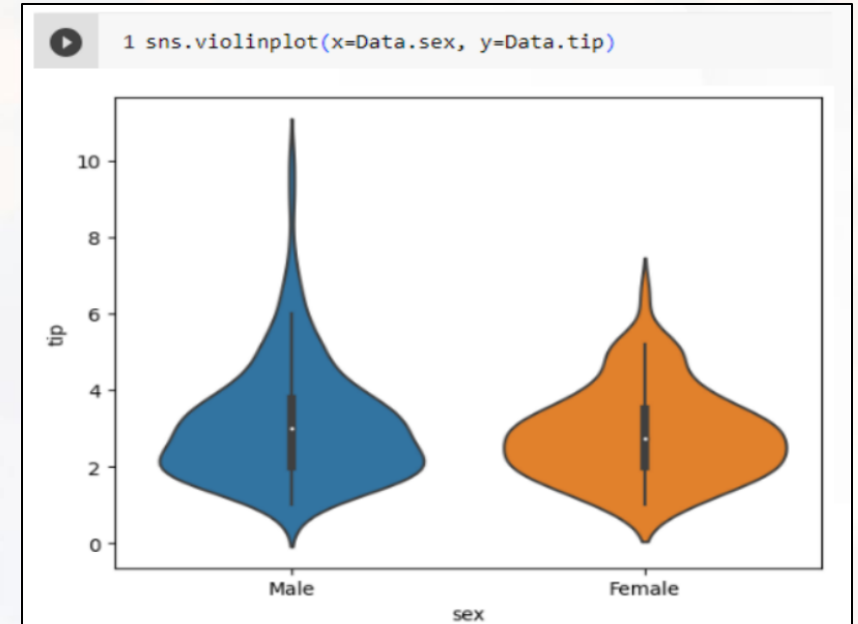




Violin / Swarm Plot


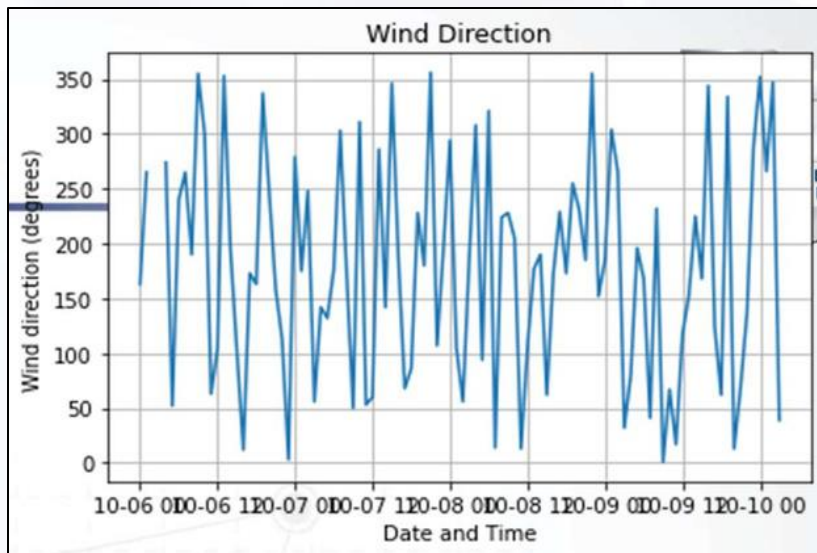
- A violin plot shows how data is spread out – whether data is evenly spread or skewed.
- To compare how data is spread across groups.

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3




Functions – Scrutinizing Datasets

- Data quality can sometimes be determined by plotting the datasets.
- Missing values or outliers can be visually identified as well as the general frequency and location.
- E.g. in the figure there is a missing value close to the beginning of the dataset and there is one entry that is an order of magnitude different from the rest. This may be a correct value or a mistake in the dataset (in many cases a domain expert would have to be consulted).



Understanding Datasets



- Information that is not always obvious when just looking at the data.
- Some features or trends can be highlighted using visualization.
- e.g., From this heat map there was, in general, poor safety in 2020 and Company 3 is a place where you would not want to work.





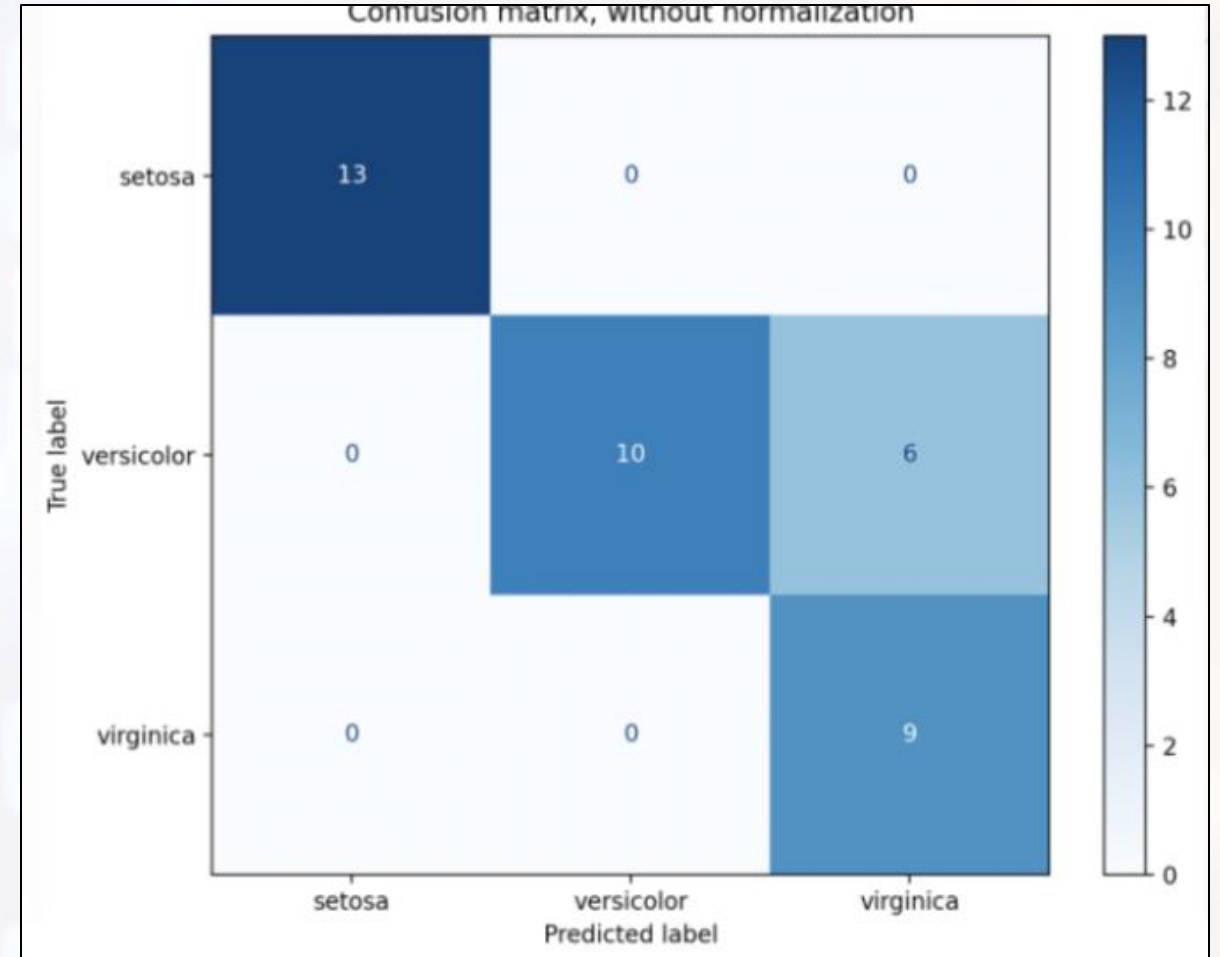
Pattern Recognition

- Clearly there is no obvious relationship between Column 1 and Column 3 in Fig 1 – which is good since it's randomized data...
- But in the second figure there is a clear linear relationship.



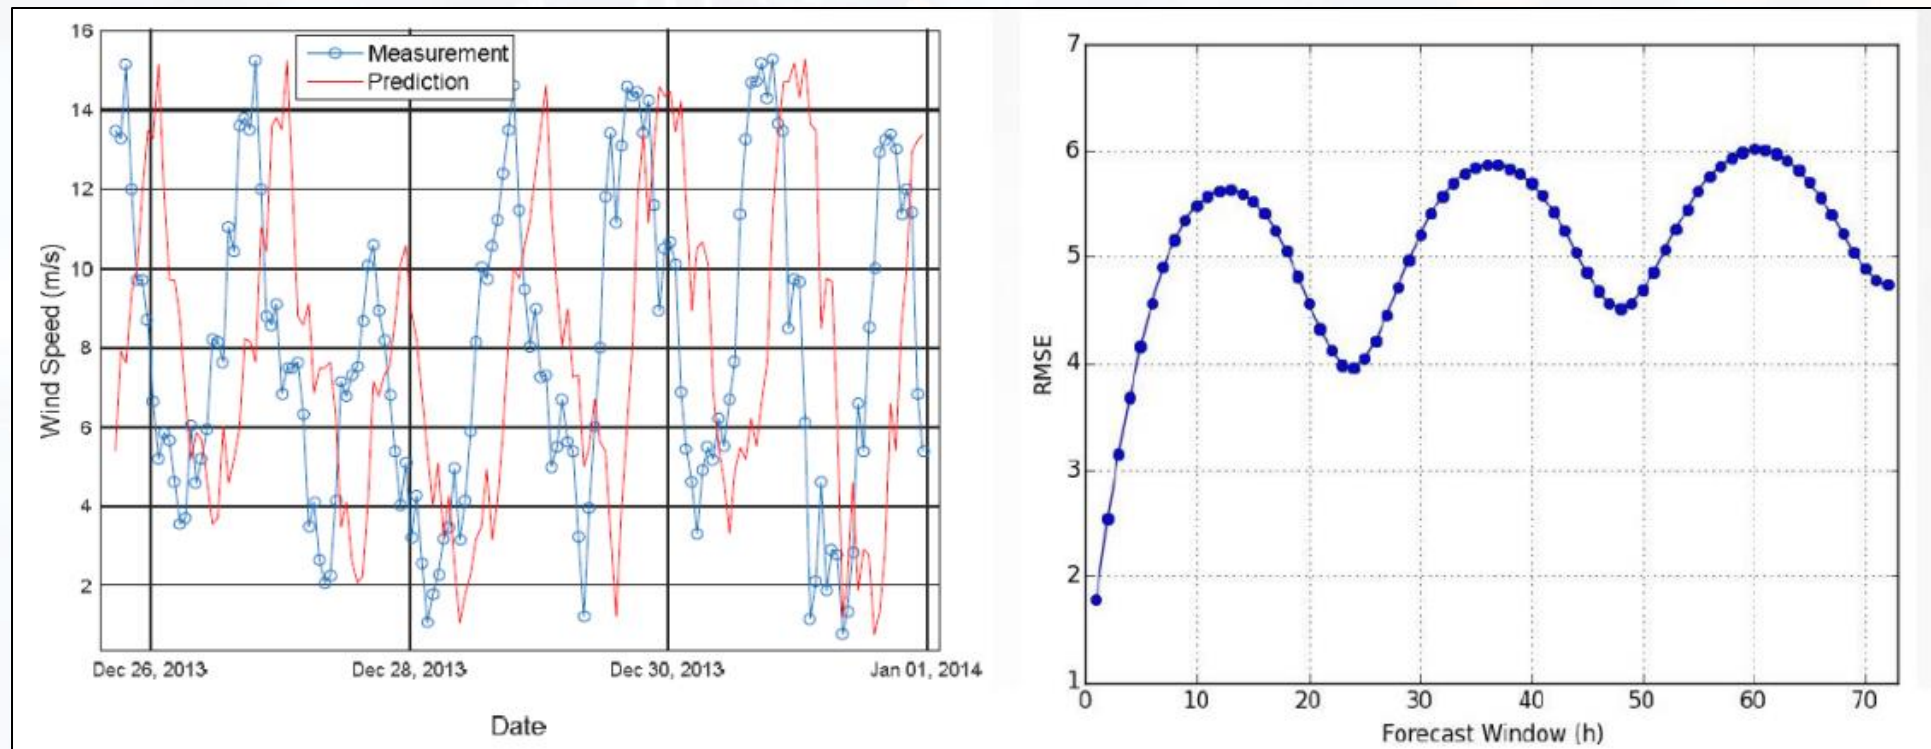
Results and Concepts Communication

- A user can better understand results from analytics and models when it is visualized.
- In cases where the actual values are of importance, the data can be displayed in an accompanying table as well.
- E.g., a confusion matrix is a very handy visual tool to help a reader understand results from a classification model.



Results and Concepts Communication

- There are also cases where it would make more sense to communicate a concept using a graph/diagram. Example:



Source: [2]



Summary

- Understanding your dataset is the first step to creating meaningful visuals
- Choosing the right type of graph depends on the message you want to show
- Well-prepared and clean data makes your visualizations more accurate and powerful
- Focus on clarity — visuals should be easy to read and support your story
- Good graphs turn raw data into insights that inform, explain, and persuade



Final Thoughts

- Never forget to label your axis, a label can change the way a graph is interpreted, and the reader or audience member will not be able to read it properly (or even yourself in 6 months' time...) if unlabelled.
- When displaying a figure or graph in a document (such as a report or thesis), always make sure to refer to it in the text and explain what is happening in the figure.
- On the other hand, in presentations, it is better to have less text. So that you would not have a written explanation of the graph on the slide, but rather explain it verbally (there are always exceptions e.g., if there is an expectation that it will be revisited at a later stage).
- When you use visualization, you should still know WHAT visualization is optimal for your problem or the concept you want to convey.
- Numerous visualization methods, as well as ways to get the graph you want, is just Google or Bing search away



References

- [1] - <https://www.sprintzeal.com/blog/fundamentals-of-data-visualization>
- [2] - <https://researchspace.csir.co.za/dspace/handle/10204/9368>

Questions???
Let's do the Data
Visualization practical

